OR Paper Review
MaGIC: Multi-modality Guided Image Completion

Bruno Sánchez Gómez

April 24, 2025

Table of Contents

Image Completion

Definition

Image completion refers to the task of filling in missing regions within an image in a visually plausible way.

- **Applications:**
  - **Inpainting:** Restoring damaged or missing parts of an image.
  - **Outpainting:** Extending the boundaries of an image.
  - **Editing:** Modifying images by adding or removing elements.
- **Approaches:**
  - **Vanilla Image Completion:** Relies solely on existing image pixels around the masked region.
  - **Guided Image Completion:** Uses external cues (e.g., text descriptions, edge maps, segmentation masks) for guidance.

Multi-modal Guided Image Completion (MaGIC) [1]

- **MaGIC:** A flexible framework for image completion guided by single or *arbitrary combinations* of modalities, such as:
  - Text
  - Canny Edge
  - Sketch
  - Segmentation
  - Depth
  - Pose
- **Architecture:** Based on pre-trained stable diffusion (SD) models with a U-Net denoiser.
- **Results:** Outperforms SOTA methods and generalizes well to various completion tasks.

Table of Contents

Components of MaGIC

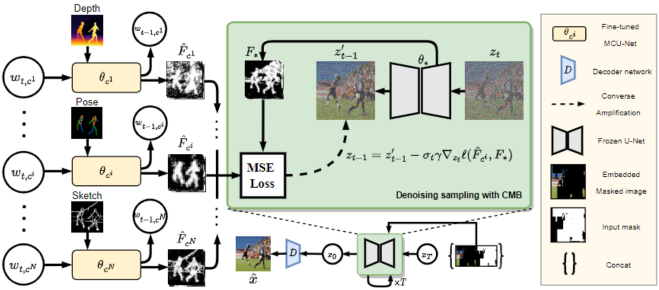

- **Modality-specific Conditional U-Net (MCU-Net):** Injects single-modal guidance into a U-Net denoiser.
- **Consistent Modality Blending (CMB):** Training-free method to blend guidance from multiple pre-trained MCU-Nets, which enables easy addition of new modalities.

Modality-specific Conditional U-Net (MCU-Net)

- The encoding network τ_c is employed to extract multi-scale guidance signals, F_c^l.
- Each F_c^l is injected to the latent in MCU-Net to obtain modality-guided features.
- To leverage pre-trained SD, the U-Net denoiser is frozen. Only the encoding network τ_c is trained to extract guidance for the frozen denoiser.
- Achieves image completion under single-modality guidance.

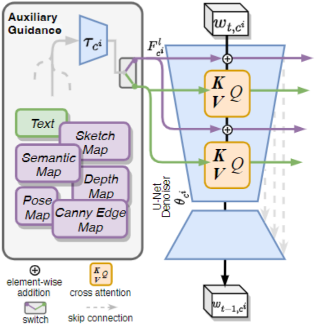

Consistent Modality Blending (CMB)

- Uses a *converse amplification strategy* [2], which enables the intermediate feature maps F_* of a original U-Net to more closely approximate the MCU-Nets' guided feature maps $\hat{F}_C$

$$\begin{cases} z_{t-1} = z'_{t-1} - \sigma_t \gamma \nabla_{z_t} \ell(\hat{F}_C, F_*) \\ \ell(\hat{F}_C, F_*) = \frac{1}{N \cdot L} \sum_{c \in C} \sum_{l=1}^{L} \left\| \hat{F}_c^l - F_*^l \right\|_2^2 \end{cases}$$

- **Properties:**
  - It is *training-free*, as it operates on already trained MCU-Nets.
  - Allows for *arbitraty combination* of available modalities.
  - Straightforward *integration of new modalities*, by simply training a new MCU-Net for them. Avoids complex joint re-training.

Quantitative Results

Method	COCO		Places2		
	FID↓	PickScore↑ / %	FID↓	U-IDS↑ / %	P-IDS↑ / %
EC (Nazeri et al., 2019) ♠	76.64	23.14	25.08	12.89	2.86
CTSDG (Guo et al., 2021) ♠	97.05	24.03	42.81	0	0
ZITS (Dong et al., 2022) ♠	61.27	28.09	18.96	18.75	7.20
Our MCU-Net†	47.70±0.29	30.79±0.10	10.74±0.07	23.83±0.30	10.18±0.48
Our MCU-Net ♡	**39.43±0.26**	**37.12±0.11**	9.09±0.04	25.34±0.29	10.64±0.46
Our MCU-Net ♣	41.91±0.20	34.96±0.17	10.27±0.06	24.21±0.24	9.93±0.38
Our MCU-Net ♠	41.15±0.27	34.94±0.06	**8.32±0.02**	**26.23±0.07**	**10.96±0.33**

Table: Comparison of using single auxiliary modality as guidance for image completion. ♠: ground truth edge map as guidance, ♡: estimated depth map as guidance, ♣: segmentation map as guidance, ↑: the higher the better, ↓: the lower the better, †: completion without any guidance.

Qualitative Results

Image examples from the paper

Table of Contents

Why did MaGIC succeed?

Where did MaGIC fail?

Future Implications [Placeholder]

- **Generalization:** MaGIC's framework can be applied to other image generation tasks, such as inpainting or super-resolution.
- **Modality Fusion:** The CMB method can be extended to fuse more complex modalities, such as audio or video.
- **Real-world Applications:** Potential applications in fields like medical imaging, autonomous driving, and augmented reality.

Thank you for your attention!

Any questions?

References I

Yongsheng Yu, Hao Wang, Tiejian Luo, Heng Fan, and Libo Zhang.
Magic: Multi-modality guided image completion.
arXiv preprint arXiv:2305.11818, 2023.

Prafulla Dhariwal and Alexander Quinn Nichol.
Diffusion models beat gans on image synthesis.
In *NeurIPS*, 2021.